

# Report

I have used 20k sentences from the training dataset to create word vectors using both Singular Value Decomposition (SVD) and Skip-gram models, with an embedding dimension of 300 for each. Additionally, for the Skip-gram with Negative Sampling (SGNS) model, I have set the number of negative samples per positive sample to 5 ( $k=5$ ).

For the downstream task, I have set certain hyperparameters as constants or fixed values to maintain consistency and ensure reproducibility.

```
input_dim = 300
hidden_dim = 128
output_dim = n_classes
n_layers = 2
bidirectional = True
n_epochs = 10
lr = 0.001
device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
```

## Hyperparameter tuning:

I trained the model using different values for the context window, specifically {1, 2, 3}, to analyze how the model's accuracies change based on the context window size.

### SVD

#### Context\_window = 1

```
Epoch 1/10, Train Loss: 1.2921832286119461, Val Loss: 1.0228583731651306
Epoch 2/10, Train Loss: 0.900489842236042, Val Loss: 0.9418131613731384
Epoch 3/10, Train Loss: 0.8792345644235611, Val Loss: 0.9685060677528381
Epoch 4/10, Train Loss: 0.7585311929583549, Val Loss: 0.7368207728862762
Epoch 5/10, Train Loss: 0.6913941984176636, Val Loss: 0.7713390529155731
Epoch 6/10, Train Loss: 0.6404095142483711, Val Loss: 0.6686382863521576
Epoch 7/10, Train Loss: 0.6113209820389748, Val Loss: 0.7125591447353363
Epoch 8/10, Train Loss: 0.573802754163742, Val Loss: 0.6239071823358536
Epoch 9/10, Train Loss: 0.5544635909199714, Val Loss: 0.6186941338777542
Epoch 10/10, Train Loss: 0.5258604573607445, Val Loss: 0.5874457001686096
```

### Metrics on Train, Validation and test set respectively

```
{'accuracy': 0.8134375, 'f1': 0.8114851738518168, 'precision': 0.8147306305298836, 'recall': 0.8134375, 'confusion_matrix': array([[3323, 341, 304, 141],
[ 154, 3727, 9, 27],
[ 281, 159, 3044, 360],
[ 316, 216, 677, 2921]])}
{'accuracy': 0.7905, 'f1': 0.7891239771458062, 'precision': 0.793769570660958, 'recall': 0.7905, 'confusion_matrix': array([[916, 66, 91, 45],
[ 88, 957, 5, 15],
[ 79, 38, 633, 79],
[ 99, 68, 181, 656]])}
{'accuracy': 0.7727631578947368, 'f1': 0.7715911211056317, 'precision': 0.7715235633255988, 'recall': 0.7727631578947368, 'confusion_matrix': array([[1475, 139, 172, 114],
[ 153, 1675, 25, 47],
[ 161, 90, 1384, 265],
[ 123, 143, 295, 1339]])}
```

## Context\_window = 2

```
Epoch 1/10, Train Loss: 1.3761390266418456, Val Loss: 1.3847731008529662
Epoch 2/10, Train Loss: 1.3867896597385407, Val Loss: 1.3805575561523438
Epoch 3/10, Train Loss: 1.139745215177536, Val Loss: 0.9213121979236603
Epoch 4/10, Train Loss: 0.8608538348674775, Val Loss: 0.8887787184715271
Epoch 5/10, Train Loss: 0.7805950139164924, Val Loss: 0.8321442592144013
Epoch 6/10, Train Loss: 0.7743128768801689, Val Loss: 0.7330533134937286
Epoch 7/10, Train Loss: 0.6961217978000641, Val Loss: 0.7920724294185638
Epoch 8/10, Train Loss: 0.6615442911982536, Val Loss: 0.9160997986793518
Epoch 9/10, Train Loss: 0.7339901869297027, Val Loss: 0.6521620433330536
Epoch 10/10, Train Loss: 0.6063337247371674, Val Loss: 0.6840354433059692
```

## Metrics on Train, Validation and test set respectively

```
{'accuracy': 0.754625, 'f1': 0.7563576759940129, 'precision': 0.7903952717761451, 'recall': 0.754625, 'confusion_matrix': array([[2942, 319, 804, 44],
[ 42, 3723, 142, 10],
[ 223, 77, 3121, 423],
[ 83, 94, 1665, 2288]])}
{'accuracy': 0.74, 'f1': 0.7469675277641789, 'precision': 0.7939804103027771, 'recall': 0.74, 'confusion_matrix': array([[820, 52, 232, 14],
[ 35, 945, 81, 4],
[ 50, 20, 672, 79],
[ 30, 31, 412, 523]])}
{'accuracy': 0.7077631578947369, 'f1': 0.7127970118227955, 'precision': 0.7555698256620181, 'recall': 0.7077631578947369, 'confusion_matrix': array([[1245, 104, 511, 40],
[ 60, 1631, 195, 14],
[ 87, 48, 1524, 241],
[ 38, 63, 820, 979]])}
```

## Context\_window=3

```
Epoch 1/10, Train Loss: 1.3838051027854283, Val Loss: 1.385075809319814
Epoch 2/10, Train Loss: 0.8773724890549978, Val Loss: 0.666205810725689
Epoch 3/10, Train Loss: 0.535577496752143, Val Loss: 0.538595265130202
Epoch 4/10, Train Loss: 0.4853864033545057, Val Loss: 0.5329739992817243
Epoch 5/10, Train Loss: 0.46062632713963586, Val Loss: 0.49787003608544667
Epoch 6/10, Train Loss: 0.43682308041801055, Val Loss: 0.46747023781140645
Epoch 7/10, Train Loss: 0.41480860993017754, Val Loss: 0.4688649616440137
Epoch 8/10, Train Loss: 0.395586516695718, Val Loss: 0.47345537998278936
Epoch 9/10, Train Loss: 0.37644791054228943, Val Loss: 0.4936177701354027
Epoch 10/10, Train Loss: 0.3625921618565917, Val Loss: 0.45585156256953874
```

## Metrics on Train, Validation and test set respectively

```
{'accuracy': 0.8740833333333333, 'f1': 0.8730298071971967, 'precision': 0.8738319111359864, 'recall': 0.8740833333333333, 'confusion_matrix': array([[21191, 1260, 924, 799],
[ 308, 23138, 68, 190],
[ 986, 769, 19651, 2354],
[ 1376, 1205, 1849, 19932]])}
{'accuracy': 0.836375, 'f1': 0.8345848586637779, 'precision': 0.8374063512850647, 'recall': 0.836375, 'confusion_matrix': array([[4931, 364, 246, 285],
[ 108, 6056, 45, 87],
[ 360, 385, 4594, 901],
[ 335, 363, 448, 4492]])}
{'accuracy': 0.8465789473684211, 'f1': 0.8451186949222905, 'precision': 0.8466953581736093, 'recall': 0.8465789473684211, 'confusion_matrix': array([[1618, 127, 75, 80],
[ 37, 1835, 8, 20],
[ 97, 84, 1476, 243],
[ 106, 122, 167, 1505]])}
```

As the size of the context window expands, we observe a corresponding rise in accuracy on the test set. This improvement is attributed to the larger window size offering a broader context, which aids in capturing extensive dependencies across the data.

## SKIP GRAM

### Context\_window=1

```
Epoch 1/10, Train Loss: 1.3215173482894897, Val Loss: 1.2432800912857056
Epoch 2/10, Train Loss: 1.002703667640686, Val Loss: 1.0120496110916137
Epoch 3/10, Train Loss: 0.6774324802160263, Val Loss: 0.7404033415317536
Epoch 4/10, Train Loss: 0.4733944005072117, Val Loss: 0.6849185242652893
Epoch 5/10, Train Loss: 0.35246706135571004, Val Loss: 0.7346596165895461
Epoch 6/10, Train Loss: 0.2672557179257274, Val Loss: 0.653699317842722
Epoch 7/10, Train Loss: 0.19069814823940395, Val Loss: 0.6979918991923332
Epoch 8/10, Train Loss: 0.14458729268424214, Val Loss: 0.783427479982376
Epoch 9/10, Train Loss: 0.10898732266761363, Val Loss: 0.8077899655103683
Epoch 10/10, Train Loss: 0.09532090242765844, Val Loss: 0.8470011223256588
```

### Metrics on Train, Validation and test set respectively

```
{'accuracy': 0.9843125, 'f1': 0.9843082249697873, 'precision': 0.9843344044780913, 'recall': 0.9843125, 'confusion_matrix': array([[4056, 22, 22, 9],
[ 32, 3873, 7, 5],
[ 18, 5, 3739, 82],
[ 6, 8, 35, 4081]])}
{'accuracy': 0.782, 'f1': 0.7831479825538817, 'precision': 0.7868997347539791, 'recall': 0.782, 'confusion_matrix': array([[912, 46, 79, 81],
[113, 849, 34, 69],
[ 92, 9, 593, 127],
[ 77, 33, 412, 724]])}
{'accuracy': 0.7759210526315789, 'f1': 0.7772030884998191, 'precision': 0.78276358252155, 'recall': 0.775921052631579, 'confusion_matrix': array([[1510, 78, 151, 161],
[ 164, 1535, 69, 132],
[ 138, 28, 1321, 413],
[ 88, 62, 219, 1531]])}
```

### Context\_Window = 2

```
Epoch 1/10, Train Loss: 1.2921832286119461, Val Loss: 1.0228583731651306
Epoch 2/10, Train Loss: 0.900489842236042, Val Loss: 0.9418131613731384
Epoch 3/10, Train Loss: 0.8792345644235611, Val Loss: 0.9685060677528381
Epoch 4/10, Train Loss: 0.7585311929583549, Val Loss: 0.7368207728862762
Epoch 5/10, Train Loss: 0.6913941984176636, Val Loss: 0.7713390529155731
Epoch 6/10, Train Loss: 0.6404095142483711, Val Loss: 0.6686382863521576
Epoch 7/10, Train Loss: 0.6113209820389748, Val Loss: 0.7125591447353363
Epoch 8/10, Train Loss: 0.573802754163742, Val Loss: 0.6239071823358536
Epoch 9/10, Train Loss: 0.5544635909199714, Val Loss: 0.6186941338777542
Epoch 10/10, Train Loss: 0.5258604573607445, Val Loss: 0.5874457001686096
```

### Metrics on Train, Validation and test set respectively

```
{'accuracy': 0.96775, 'f1': 0.967638912926545, 'precision': 0.9679580195393938, 'recall': 0.96775, 'confusion_matrix': array([[3940, 61, 73, 35],
[ 29, 3850, 9, 29],
[ 12, 22, 3739, 71],
[ 26, 26, 123, 3955]])}
{'accuracy': 0.765, 'f1': 0.7668977461816581, 'precision': 0.7722147486338684, 'recall': 0.765, 'confusion_matrix': array([[824, 73, 128, 93],
[ 66, 883, 54, 62],
[ 51, 19, 629, 122],
[ 72, 38, 162, 724]])}
{'accuracy': 0.7610526315789473, 'f1': 0.7630590213785383, 'precision': 0.7694956191176988, 'recall': 0.7610526315789473, 'confusion_matrix': array([[1382, 118, 232, 168],
[ 106, 1524, 105, 165],
[ 67, 50, 1421, 362],
[ 92, 70, 281, 1457]])}
```

### Context\_window = 3

```
Epoch 1/10, Train Loss: 1.288504887342453, Val Loss: 1.1797360997200013
Epoch 2/10, Train Loss: 1.0649072164297104, Val Loss: 1.0468833956718444
Epoch 3/10, Train Loss: 0.7982144811749459, Val Loss: 0.8037800433635711
Epoch 4/10, Train Loss: 0.6125875627994537, Val Loss: 0.7095185635089875
Epoch 5/10, Train Loss: 0.4768124467283487, Val Loss: 0.7473996315002441
Epoch 6/10, Train Loss: 0.38755657204985616, Val Loss: 0.6915469779968262
Epoch 7/10, Train Loss: 0.3194863688647747, Val Loss: 0.6924065791368484
Epoch 8/10, Train Loss: 0.24967132564261554, Val Loss: 0.7260419960021973
Epoch 9/10, Train Loss: 0.2030205830335617, Val Loss: 0.7169653385877609
Epoch 10/10, Train Loss: 0.18421075877919793, Val Loss: 0.7354234385490418
```

## Metrics on Train, Validation and test set respectively

```
{'accuracy': 0.9648125, 'f1': 0.9647853072138769, 'precision': 0.9651259161549824, 'recall': 0.9648125, 'confusion_matrix': array([[3862, 61, 138, 48],
       [ 12, 3889, 7, 9],
       [ 28, 19, 3671, 126],
       [ 17, 15, 83, 4015]])}
{'accuracy': 0.78675, 'f1': 0.7878349637540384, 'precision': 0.798103726648105, 'recall': 0.78675, 'confusion_matrix': array([[857, 69, 94, 98],
       [ 77, 906, 46, 36],
       [ 60, 18, 619, 124],
       [ 57, 37, 137, 765]])}
{'accuracy': 0.776578947368421, 'f1': 0.7781953212284034, 'precision': 0.7819345236554393, 'recall': 0.776578947368421, 'confusion_matrix': array([[1468, 76, 166, 190],
       [ 98, 1588, 106, 116],
       [ 121, 43, 1382, 354],
       [ 105, 68, 255, 1472]])}
```

However here as context\_window size changes the accuracy doesn't seem to be affecting much. This might be since we are training the word embeddings on a small dataset, and the test accuracy is already high. Enlarging the context window in a skip-gram model may not consistently enhance accuracy, often due to increased semantic drift, data sparsity, heightened computational complexity, loss of local context, and the risk of overfitting.

## CONCLUSION:

We observe that skip-gram model (with negative sampling) performs better than svd in all metrics.